

The predation problem, reciprocity, and the limits of consequentialism

Abstract: Our article examines the predation problem in animal ethics as a challenge to the combination of Singer's principle of equal consideration of interests and welfare consequentialism. Equal consideration requires that the suffering and enjoyment of non-human animals be counted impartially, while welfare consequentialism directs moral evaluation toward net welfare balance. This combination, however, appears to yield counterintuitive and even monstrous conclusions about certain cases of predation. We argue that this difficulty cannot be resolved by empirical considerations, an appeal to a side constraints, rights-based protections, or a deontological floor, or an appeal to rule-based versions of consequentialism. Instead, the wrongness of predation in the relevant cases is best understood through the structure of reciprocal moral justification: moral evaluation belongs to a public, rule-governed space of reasons. What emerges from our analysis is a philosophically rich, two-tiered structure that maintains welfare and reciprocity in productive tension, as our symmetry-in-the-space-of-reasons response is combined with Singer's principle of equal consideration of interests and welfare consequentialism.

Keywords: Predation problem; animal ethics; Singer; welfare consequentialism; reciprocity; space of reasons; normativity; value

§1 Introduction

Peter Singer’s (1989; 1977, 1979, 2002, 2023) principle of equal consideration of interests concerns all individuals with interests – whether simple or complex – and requires us human beings to consider equally all the interests of those affected by our actions. If all sentient beings have interests, then the principle of equal consideration of interests requires that we give equal consideration to their interests, irrespective of whether they are members of our species or of another species.¹ This principle, rejecting the claim that an animal’s suffering matters less simply because it is a non-human animal rather than a human that is suffering, became a field-shaping, foundational idea in modern animal ethics. Combined with welfare consequentialism, Singer’s principle of equal consideration of interests grounds powerful objections to factory farming, animal testing for non-essential (for instance, cosmetic) products, culling practices, selective breeding, certain forms of entertainment involving the use of animals, intensive fish farming, neglect in pet ownership, and other practices that impose avoidable suffering on non-human animals.

At the same time, the combination of Singer’s principle and welfare consequentialism appears to generate a puzzle about predation. The predation problem has already received much philosophical attention (Clark, 1979; Ebert and Machan, 2012; Keulartz, 2016; Ng, 1995; Ritchie, 1895).² For the

¹ Failure to meet this requirement incurs the charge of speciesism, an ethically indefensible form of discrimination against beings on the basis of their membership of a species other than our own.

² The predation problem can be framed through a series of interconnected questions. Should wild nature be left alone because predation is integral to ecological balance, or is suffering in nature somehow morally significant? Can predator-caused suffering be bad even if the predators themselves are neither evil nor morally blameworthy? If such suffering matters (morally speaking), do we humans have reasons – or perhaps even

sake of argument, our use of the term ‘predation’ is intended to cover animal-on-animal (lion versus gazelle), human-on-animal (hunting, factory farming), human-on-human (cannibalism as a special case), and alien-on-human (intelligent or superintelligent predators) predation.³ In at least some cases of predation (animal-on-animal, human-on-animal, human-on-human, and alien-on-human), a predator’s gains in pleasure, survival, and relief from hunger or suffering might outweigh the prey’s losses, especially when the suffering of the prey is brief. If moral evaluation tracks the net balance of affected interests in a manner that is both impartial and consistent with Singer’s principle of equal consideration of interests and welfare consequentialism, then these cases will appear favourable under a welfare-based calculus. Nonetheless, few would conclude that predation is therefore morally permissible, let alone obligatory under these circumstances. This creates a three-way tension: we cannot simultaneously hold that (a) non-human animal interests count impartially (Singer’s principle); (b) moral rightness is determined entirely by net welfare considerations (welfare consequentialism); and (c) we can retain moral reservations about predation, even when the welfare balance appears favourable.

As we shall argue, the central problem is that an impartial, welfare-based framework cannot explain how we can retain moral reservations about predation, even when the welfare balance appears favourable. It tells us that animal suffering counts, but not why we can retain reservations about predation in certain instances, even when the arithmetic seems favourable. Something beyond Singer’s principle of equal consideration of interests and welfare consequentialism is doing the normative work, and we identify the

obligations – to intervene, whether through population control, habitat management, or more speculative possibilities such as genetically engineering carnivores into herbivores?

³ More generally, we use ‘predation’ to cover all cases where one sentient being treats another as consumable resources, including human cannibalism as a limiting case.

missing element as a structural constraint grounded in reciprocal rational agency. In particular, certain cases of predation place the predator outside the symmetry of mutual justification required for moral evaluation. Once this constraint is recognized, equal consideration and welfare consequentialism can be preserved, while allowing us to rule out certain cases of predation on principled grounds.

§2 The argument

An argument, taking the form of a *reductio ad absurdum*, can be constructed in support of our claim concerning the inadequacy of equal consideration plus welfare aggregation:

P1: Singer's principle of equal consideration of interests requires that the suffering and enjoyment of non-human animals be counted morally, not discounted or ignored because these animals are non-human.

P2: If suffering and enjoyment are counted impartially, then the moral evaluation of each action must track the overall balance of affected interests.

P3: In at least some cases of predation, the predator's gains in pleasure, survival, and relief from hunger or suffering appear to outweigh the prey's losses.

P4: If moral rightness is completely determined by the overall balance of affected interests, then some predation would be morally permissible, perhaps even obligatory.

P5: However, predation is not morally permissible in these cases. The fact that an action is welfare-maximizing does not, by itself,

make it an acceptable form of treatment of sentient beings.

C: $\therefore$ Equal consideration plus welfare consequentialism are insufficient to ground a stable, anti-predation ethic.

Given the form of the reductio that our argument takes, the question becomes: which premise should be rejected and on what grounds? Alternatively, the three-way tension between (a) non-human animal interests counting impartially; (b) moral rightness being fixed entirely by the net balance of welfare; and (c) predation remaining morally wrong, even when the welfare balance appears favourable, can be restated in the form of a trilemma:

(Horn 1 or impartiality) Non-human animal interests counting impartially and morally;

(Horn 2 or welfare consequentialism) Moral rightness is solely determined by the net balance of welfare;

(Horn 3 or anti-predation constraint) Some predation remains wrong, even when the net welfare balance seems favourable.

Our argument implies that the three claims associated with Horns 1, 2, and 3 cannot be simultaneously true. If we keep Horns 1 and 2 and ditch Horn 3, then we must concede that predation will be justified whenever the net welfare balance seems favourable. Such a move could generate a cascade of deeply counterintuitive and potentially monstrous implications. If we keep Horns 2 and 3 and discard Horn 1, then we might have to reintroduce some form of morally relevant distinction between humans and other animals or between different classes of animals (predators versus prey). Such a move would look uncomfortably close to the speciesism that Singer's principle was designed to reject. If we keep Horns 1 and 3 and ditch Horn 2, then we might have to supplement or replace welfare consequentialism

with a framework that includes side-constraints, agent-relative permissions, or rights-based protections that hold even when violating these constraints or protections would result in a net welfare gain. The cost of this move involves having to explain what grounds these additional constraints, permissions, or protections, if not welfare itself, and how these constraints, permissions, or protections would apply to non-moral agents such as predatory animals.

We are reluctant to abandon the combination of Singer's principle of equal consideration of interests (Horn 1) and welfare consequentialism (Horn 2). Together, Horns 1 and 2 have provided the theoretical foundations for the development of the bulk of modern animal ethics over the past half-century, and much of what is most defensible in this tradition depends on retaining some version of these two horns. What is needed, then is not a rejection of the theoretical foundations associated with Horns 1 and 2 but rather a supplement to these foundations – some account, in other words, of why Horn 3 retains its force even within a broadly impartial and welfare-based framework, and of what theoretical resources this framework might draw on or at least acknowledge in order to accommodate it.

§3 Standard responses to the predation problem

One standard response to the predation problem involves a denial of P3 of our argument (§ 2): perhaps predation doesn't actually maximize the net welfare when all the consequences are taken into account (for instance, the effects on the ecosystem). This response is problematic for at least two reasons. Firstly, it is a contingent response. Even if predation usually produces suboptimal net welfare, would predation be permissible – at least in principle – when the moral calculus genuinely favoured it? If our answer is 'yes', then we have conceded that predation is in principle permissible under certain welfare

conditions. Most would find this concession, which Horn 3 is designed to resist, deeply counterintuitive and troubling. If our answer is ‘no’, then we have implicitly acknowledged that something beyond our combination of Singer’s principle (P1) and welfare consequentialism (P2) must be doing the moral work. Secondly, this response flounders in structurally similar cases where the welfare calculus would clearly favour the predatory practice:

(Repugnant predation) Suppose that a trillion tiny yet intelligent aliens each derive a miniscule yet non-zero unit of pleasure from feasting on a single human subject. This human is not killed outright, but rather sustained in a state of chronic suffering across her entire lifespan, so the feast endures. While each alien’s pleasure is negligible in isolation, the cumulative welfare gain, aggregated across a trillion consumers, would vastly exceed the welfare loss of the single human subject. The moral calculus, applied impartially, would return the verdict that this arrangement is not only morally permissible but also obligatory;

(Hedonivores) Suppose that a superintelligent alien species – call them the Hedonivores – experienced pleasure a trillion times more intensely than humans experience either pain or pleasure. Their neurobiological substrate is vastly richer and more complex, and their hedonic capacities incomparably deeper, such that a single moment of gustatory satisfaction generates more welfare, by our moral calculus’s own lights, than an entire human lifetime of flourishing. These Hedonivores take no special pleasure in our suffering and in fact are indifferent to it. Possessing superior civilization, culture, and social relationships, they are, by every single measure, more richly conscious and cognitively sophisticated

than we are. Whenever one of these Hedonivores consumes a human being, the welfare gain is astronomical and the welfare loss, though real, is negligible by comparison. Without species discount and applied impartially, the moral calculus does not merely permit these Hedonivores to consume us – it appears to require that we facilitate and enable this arrangement and offer ourselves up.

Even if we grant every number, the conclusions that the single human subject ought to be sacrificed for the aggregate good, and that we humans ought to be consumed by these Hedonivores, are not merely counterintuitive but monstrous. Crucially, the monstrousness of these conclusions is not a function of the arithmetic. These conclusions do not become acceptable after the sums are checked and confirmed: they seem wrong in a way that is entirely independent of whether the welfare calculus has been correctly applied. The upshot is that wrongness is located elsewhere – where the welfare calculus, however scrupulously applied, has no dominion.

Another standard response to the predation problem involves the claim that something is constraining the verdict in favour of Horn 3: a side constraint, a rights-based protection, or a deontological floor, that is not itself grounded in welfare maximization. For instance, Tom Regan (1985) and other philosophers have argued that at least some non-human animals have rights that constrain welfare maximization. Furthermore, it may be asserted that predation violates these rights, regardless of the net welfare balance. This standard response avoids the predation problem, albeit at a significant cost. Firstly, it jettisons the consequentialist framework completely, instead of attempting to resolve the internal tension within it. Secondly, it leaves unexplained where the additional side constraints, rights, or deontological floors come from and why they have the specific structure that they do. The

appeal to side constraints, rights-based protections, or even a deontological floor appears ad hoc, unless this appeal is grounded in a more fundamental account of what makes exploitation morally wrong.

Yet another standard response to the predation problem might attempt to resolve the internal tension within consequentialism by arguing that rules that prohibit predation in the repugnant predation, Hedonivores, and other structurally similar cases tend to give rise to better overall consequences than act-based consequentialism. This response, however, simply relocates the question: why do anti-predation constraints and rules give rise to better consequences? If the rule-consequentialist answer is that predation causes more suffering than act consequentialists typically calculate, we have already pointed out that these predatory actions seem wrong in a way that is independent of our getting the mathematics right. If the rule-consequentialist answer is that predation has some morally wrong features that are independent of welfare-based considerations, then we would have conceded the insufficiency of welfare consequentialism. All these standard responses – a failure to maximize net welfare when all the relevant consequences are taken into account; an appeal to side constraints, rights, or deontological floors; an appeal to rule-based forms of consequentialism – share an inability to explain why predation is wrong in a manner that integrates both impartial and welfare-based considerations. Therefore, we need a positive account of what makes treating sentient beings as consumable resources morally impermissible in the repugnant predation, Hedonivore, and other structurally similar cases, even when the welfare balance seems favourable to the would-be predators.

§4 Our response

Our response begins with the simple observation that moral evaluation does not occur in a vacuum: it already presupposes a structure of moral justification among rational agents, such that asking ‘What ought to be done?’ is already to participate in the practice of demanding and providing reasons that others can recognize. This observation can be situated within a broader Wittgenstein-Sellars-Brandom lineage, according to which normativity – and by implication, moral evaluation – is socially instituted, rule-governed, and essentially public, rather than grounded in private mental states. Wittgenstein’s (1953, §§201–202) reflections on rule-following remind us that the correct application of a rule cannot be fixed by individual interpretation alone, but rather through participation in a shared practice. Sellars (1968, 1980) develops this Wittgensteinian insight by construing moral principles as we-intentions or shared commitments, articulated from the standpoint of a collective ‘we’. For instance, ‘I shall help this injured animal’ (a subjective preference) becomes ‘I shall_{we} help this injured animal’ (an intersubjectively intelligible commitment).⁴ Brandom (1994) then systematized this insight through his account of a space of reasons, a normative structure implicit in the practice of giving and asking for reasons. In this space of reasons, an abstraction from the rule-governed social practice, agents undertake and attribute commitments and entitlements and rationality consists in a communal, score-keeping game of giving and asking for reasons, not inner episodes of thought. The roots of our observation can equally be traced to Kant’s (1785) categorical imperative, Rawls’s (1971) veil of ignorance, Scanlon’s (1998) notion of reasonable rejectability, and Habermas’s (1991, 1993a,b) discourse ethics. At the same time, we need not commit ourselves to any

⁴ When the grammatical subject in a we-intention is not ‘we’, Sellars indicates its status as a we-intention by using ‘we’ as a subscript.

particular version of these views. The core claim in our response is simply that normativity arises from the structure of reciprocal justification.

Next, our response distinguishes between value and normativity. An account of value is needed if we wish to make sense of what we are giving reasons about or what might make something worth justifying or condemning. Consciousness is the ontological ground for value, since there is something it is like to suffer or experience pain or pleasure (Nagel, 1974). Pleasure, suffering, flourishing, and pain are not inferred through reasoning but immediately given as experienced aspects of phenomenal reality. In addition, these hedonic states matter from the inside and possess intrinsic value for each consciousness. By contrast, a universe filled with lifeless, unconscious matter would be indifferent, containing neither triumph nor tragedy, value nor disvalue.⁵

If consciousness grounds value, then reciprocity grounds normativity. Although it is only my own suffering that I can experience first-hand, if I claim that my suffering matters more simply because it is mine, then I offer no reason that other rational agents in a reciprocal structure could recognize as valid. Shared rationality therefore demands a standpoint that abstracts away from the contingencies and arbitrariness of private experience, in favour of the public, rule-governed space of reasons. The deeper logic of shared rationality can be found in Rawls's (1971) veil of ignorance, Sidgwick's (1874) axiom of rational benevolence, and Hare's (1963) notion of universalizability. We now have the resources available to explain what is wrong with predation, especially in the repugnant predation and Hedonivores scenarios (§ 3). When

⁵ As far as we can tell, the observable universe spent roughly its first nine billion years after the Big Bang in an abiotic state, containing neither life nor consciousness. Without consciousness and phenomenal experience as the ground for value, then not a single instant of those nine billion years would have contained even the slightest trace of value.

a predator treats a prey as a consumable resource, it places itself outside the reciprocal justificatory symmetry associated with the space of reasons.

In our repugnant predation and Hedonivores scenarios, there is something self-undermining about a rational agent – any of the trillion tiny yet intelligent aliens in the former scenario and the Hedonivore in the latter – invoking the moral calculus to justify predation and operating within the space of reasons, while denying moral standing to its prey. Roughly, this agent is claiming that the prey’s suffering counts – that is why the welfare calculus is being used in the first place. At the same time, the predator-prey relationship, in which the prey is treated as a consumable resource, somehow exempts this agent from treating the prey as a participant in justification. The space of reasons carries with it reciprocal obligations of justification and reason-sharing that the predator-prey relationship simultaneously violates. At least some cases of predation are morally wrong, not merely because they cause suffering, but because they violate the constitutive conditions of moral evaluation: a rational agent cannot coherently engage in predation, while claiming to participate in reciprocal moral justification.

By contrast, lions preying on gazelles are not participants in the space of reasons. They cannot demand or give reasons, undertake commitments, or be held to justificatory account. Since they fall outside the normative framework altogether, they are morally blameless in their predatory actions.⁶ The self-undermining charge targets rational agents who invoke moral authority, while simultaneously denying moral standing to those whom they harm. There is a clear analogy here between the predator-prey relationship and the exploiter-

⁶ The lion-gazelle case does raise a separate question of whether we, as rational agents, have a moral obligation to intervene in predation in the wild. However, we are making a different argument here about the moral permissibility of certain cases of predation (for instance, predators as rational agents consuming their prey).

exploited relationship. The space of reasons does not require the prey to be a participant in the practice of demanding or giving reasons – only that they could be a subject of this practice. Rational agents – including the predators in the repugnant predation, Hedonivores, human-on-animal, and human-on-human scenarios – who invoke the welfare calculus are obliged to justify their actions to a normative community of other rational agents. This community can, however, include beings who cannot speak for themselves, since their interests can be represented by other rational agents. The prey need not be capable of demanding reasons for being treated in a certain fashion: it will be sufficient that reasons must be given about how these beings are treated.

When our symmetry-in-the-space-of-reasons response is combined with Singer’s principle of equal consideration of interests and welfare consequentialism, we end up with a two-tiered view:

- (Tier 1 or welfare) All sentient beings deserve equal consideration and moral rightness is completely determined by the net balance of welfare (Horns 1 and 2);
- (Tier 2 or reciprocity) Rational agents operate in a public, rule-governed space of reasons and owe each other additional duties of reciprocal justification (Horn 3).

In the scenario of the lion eating the gazelle, the gazelle is sentient and has interests in avoiding pain and suffering and continuing its existence. These interests also count equally in the welfare calculus. While the lion’s predatory action causes negative welfare (the suffering and death of the gazelle), it also produces positive welfare (the lion’s survival). Under welfare consequentialism, this is a tragic welfare tradeoff (Tier 1). At the same time, as the lion is not a rational agent operating in the space of reasons, it cannot respond to normative claims, offer justification, or recognize

duties. Therefore, no reciprocity obligations apply and we cannot hold the lion morally blameworthy or demand justification from it ($\neg$ Tier 2). The lion's predation is therefore both welfare-relevant (Tier 1) and not a moral wrong whose blameworthiness can be attributable to an agent ($\neg$ Tier 2). In the scenario of humans raising and killing cattle for beef, cattle are as sentient as gazelles and their interests in avoiding pain, suffering, and death count equally under Singer's principle of equal consideration. Furthermore, factory farming produces enormous net negative welfare (for instance, suffering during confinement, transport, and slaughter) that far outweighs the welfare gains (for instance, human nutrition, gustatory pleasure, and so on) (Tier 1). At the same time, humans are rational agents in the space of reasons who can recognize normative claims, and demand and offer reasons (Tier 2). This generates at least two distinct obligations: the obligation not to cause unnecessary suffering (owed to the cattle at Tier 1) and the obligation to publicly justify practices that might impose prohibitively high welfare costs (at Tier 2).

In the human-on-human (cannibalism) scenario, both tiers are activated simultaneously and convergently. At Tier 1, the human victim has interests in survival, avoiding pain, and having future experiences. The welfare loss to the victim, all other things being equal, will outweigh the predator's gain, especially given other available food alternatives. At Tier 2, the victim is a rational agent in the space of reasons, owed not merely welfare-based considerations but also reciprocal justification. Killing humans for food and treating them as a mere resource to be consumed results in their standing as a rational agent being bypassed. Cannibalism is therefore doubly wrong: Tier 1 and 2 violations are compounded, producing the strongest possible condemnation. In certain outlier instances, however, Tiers 1 and 2 might not be mutually reinforcing. When the victim is already dead (for instance, a corpse

on a lifeboat), there will not be any interests to speak of or provide equal consideration for. Under strict welfare consequentialism, no welfare-based considerations will accrue to the corpse. At the same time, the survivors' welfare interests are real and urgent. The welfare calculus appears to favour consumption. At Tier 2, since the deceased can no longer participate in the space of reasons, no reciprocity is owed to them.

In the repugnant predation scenario, the predation appears not merely morally permissible but even obligatory at Tier 1. This scenario demonstrates how naive welfare aggregation has the following pathological attributes:

(Unbounded replaceability) Any finite loss in welfare by a single individual can in principle be outweighed by sufficiently many small welfare gains distributed across a sufficiently large number of beings. This turns the prey into a welfare battery;

(Weaponization of impartiality) Singer's principle of equal consideration, intended to expand the moral circle, ends up being turned against this welfare battery;

(Duration as a multiplier) The human victim is not immediately killed but sustained in suffering. This scenario parameter is a deliberate choice that allows for the aggregate alien welfare to be maximized over time. Prolonged suffering of the welfare battery to extend aggregate pleasure is not merely morally permissible but obligatory.

Taken on its own, the Tier 1 verdict will be in favour of the single human subject being turned into a welfare battery. At the same time, this human subject is a rational agent in the space of reasons. At Tier 2, rational agents should not be treated as mere welfare batteries without any further justification that they could in principle accept. In addition, the trillion-alien

arrangement results in the victim's rational standing being overridden by the welfare calculus. This human subject could reasonably reject any principle that licenses her being used in the manner prescribed by the repugnant predation scenario, regardless of the aggregate benefits that might accrue. The Tier 2 verdict, by contrast, will be against the trillion-alien, repugnant predation arrangement, independently of the arithmetic.

In the Hedonivores scenario, the Hedonivores are by stipulation more richly conscious, cognitively sophisticated, and culturally advanced. In addition, they take no special pleasure in human suffering. Given how these Hedonivores are constituted, the welfare gains that are generated per moment of consumption will dwarf an entire human lifetime of flourishing. Singer's principle of equal consideration asserts that equal interests count equally. However, the Hedonivores' interests are, by stipulation, incomparably greater in magnitude. The Hedonivores case generates three pathological attributes under Tier 1:

(Consciousness gradient) If the magnitude of welfare tracks the degree of richness or sophistication of consciousness, and equal consideration means equal weight per unit of welfare rather than equal weight per being, then beings with vastly superior hedonic capacities will dominate welfare-based calculations. The principle of equal consideration of interests will become, despite Singer's best intentions, a principle under which the interests of the less richly conscious are subordinated by the interests of the more richly conscious;

(Incommensurability across substrates) The Hedonivore case stipulates that the Hedonivores experience pleasure a trillion times more intensely than human beings. This stipulation is doing much of the philosophical work. If, however, we cannot compare

welfare across distinct neurobiological substrates in any meaningful fashion, then the welfare calculus becomes infeasible or impracticable;

(Unbounded replaceability) In the repugnant predation case, the number of predators overwhelms the welfare of the individual victim. In the Hedonivores scenario, the intensity of the predator's welfare overwhelms the welfare of the victim. The welfare consequentialist framework appears to be vulnerable along multiple dimensions of aggregative structure.

At Tier 2, we expect Hedonivores to be fully subject to reciprocity obligations: they owe human beings a justification for their predatory practices. The question in the Hedonivores scenario becomes: can the consumption of human beings be justified within a shared, public, rule-governed space of reasons that includes both our species and the Hedonivores'?

§5 Conclusion

Singer's principle of equal consideration of interests and welfare consequentialism have together provided the theoretical foundations for the development of the bulk of modern animal ethics over the past half-century. At the same time, the predation problem reveals a deep, underlying tension within welfare-based approaches to animal ethics. In particular, if we maintain Singer's principle of equal consideration of interests in conjunction with the welfare consequentialist's belief that moral rightness is determined solely by the net balance of welfare, then we seem forced into troubling conclusions about certain cases of predation in which the consumption of prey would be justified by aggregate welfare gains. This result is not merely counterintu-

itive: it suggests that the equal consideration of interests and welfare-based considerations do not, on their own, capture what is morally at stake when a rational agent determines whether to treat another being as a consumable resource.

As we have argued, the mistake involves supposing that moral judgment is nothing over and above the correct application of the welfare calculus, detached from the structure of reciprocal justification. Whereas phenomenal consciousness grounds value, reciprocity grounds normativity: moral evaluation belongs to a public, rule-governed space of reasons where rational agents are answerable to one another and must be able to offer reasons that others can recognize as valid or – at least – will not be able to reasonably reject. In this setting, predation by rational agents is not wrong simply because it causes suffering, but because it violates the symmetry of justification that moral agency presupposes. We do not need to abandon either the equal consideration of interests or welfare-based considerations. Instead, we simply need to supplement these commitments with a further normative condition: moral agency must be governed by reciprocity, not merely welfare aggregation. Once this normative condition is recognized, the moral wrongness of predation becomes intelligible, without our having to either revert to speciesism or give up an impartial concern for suffering in all its guises. What our analysis yields, then, is a more adequate moral framework, one that can respect the interests of sentient beings while explaining why certain forms of predation and exploitation remain morally impermissible, even when the moral arithmetic might appear favourable.

More generally, when our symmetry-in-the-space-of-reasons response is combined with Singer’s principle of equal considerations of interest and welfare consequentialism, the result is a two-tiered view: welfare at Tier 1 and reciprocity at Tier 2. Our two-tiered view proves, as it turns out, to be a

remarkably generative structure for navigating our rich gamut of predation scenarios. The lion-gazelle case establishes the baseline: welfare-based considerations are relevant (Tier 1), but the absence of rational agency means no moral wrong can be attributable ($\neg$ -Tier 2). In the human-cattle scenario, rational agents incur obligations both to minimize suffering (Tier 1) and publicly justify practices that impose welfare costs on sentient beings (Tier 2). Human cannibalism compounds violations at Tiers 1 and 2. The repugnant predation case, in turn, stress-tests our framework, revealing that Tier 1, taken on its own, is vulnerable to certain pathologies, making Tier 2 not merely an additive supplement but a lexical side-constraint. The Hedonivore case arguably exposes the deepest tension within our two-tiered view: if welfare comparisons across distinct neurobiological substrates can be coherently made, then it is possible that deeply counterintuitive Tier 1 verdicts will be generated that even a robust Tier 2 lexical constraint might struggle to contain. What emerges from our analysis is not some tidy and straightforward decision procedure but rather a philosophically rich structure that maintains welfare (Tier 1) and reciprocity (Tier 2) in productive tension, as our symmetry-in-the-space-of-reasons response is combined with Singer's principle of equal consideration of interests and welfare consequentialism. In certain instances of predation (the repugnant predation case), welfare consequentialism may require the corrective force of reciprocal justification. In other instances (the Hedonivores case), even this corrective force might turn out to be insufficient.

References

Brandom, Robert (1994). *Making It Explicit: Reasoning, Representing, & Discursive Commitment*. Cambridge: Harvard University Press.

- Clark, Stephen R. L. (1979). “The Rights of Wild Things”. In: *Inquiry: An Interdisciplinary Journal of Philosophy* 22.1-4, pp. 171–188. DOI: 10.1080/00201747908601871.
- Ebert, Rainer and Tibor R. Machan (2012). “Innocent Threats & the Moral Problem of Carnivorous Animals”. In: *Journal of Applied Philosophy* 29.2, pp. 146–159. DOI: 10.1111/j.1468-5930.2012.00561.x.
- Habermas, Jürgen (1991). *The Theory of Communicative Action: Reason & the Rationalization of Society*. Polity.
- (1993a). *Justification & Application: Remarks on Discourse Ethics*. Polity.
- (1993b). *Moral Consciousness & Communicative Action*. Polity.
- Hare, Richard Mervyn (1963). *Freedom & Reason*. Oxford, Clarendon Press.
- Kant, Immanuel (1785). *Groundwork for the Metaphysics of Morals*. Ed. by Thomas E. Hill and Arnulf Zweig. New York: Oxford University Press.
- Keulartz, Jozef (2016). “Should the Lion Eat Straw Like the Ox? Animal Ethics & the Predation Problem”. In: *Journal of Agricultural & Environmental Ethics* 29.5, pp. 813–834. DOI: 10.1007/s10806-016-9637-4.
- Nagel, Thomas (1974). “What is It Like to Be a Bat?” In: *Philosophical Review* 83.4, pp. 435–50. DOI: 10.2307/2183914.
- Ng, Yew-Kwang (1995). “Towards Welfare Biology: Evolutionary Economics of Animal Consciousness & Suffering”. In: *Biology & Philosophy* 10.3, pp. 255–285. DOI: 10.1007/bf00852469.
- Rawls, John (1971). *A Theory of Justice*. Ed. by Steven M. Cahn. Oxford, Harvard University Press.
- Regan, Tom (1985). *The Case for Animal Rights*. University of California Press.
- Regan, Tom and Peter Singer, eds. (1989). *Animal Rights & Human Obligations*. Cambridge University Press.
- Ritchie, David G. (1895). *Natural Rights: A Criticism of Some Political & Ethical Conceptions*. Routledge.

- Scanlon, Thomas, ed. (1998). *What We Owe to Each Other*. Cambridge: Harvard University Press.
- Sellars, Wilfrid (1968). *Science & Metaphysics: Variations on Kantian Themes*. New York, Humanities P.
- (1980). “On Reasoning About Values”. In: *American Philosophical Quarterly* 17.2, pp. 81–101.
- Sidgwick, Henry (1874). *The Methods of Ethics*. Ed. by Emily Elizabeth Constance Jones. Bristol, U.K.: Cambridge University Press.
- Singer, Peter, ed. (1977). *Animal Liberation*. Avon Books.
- (1979). *Practical Ethics*. Ed. by Susan J. Armstrong and Richard George Botzler. New York: Cambridge University Press.
- (2002). “Ethics across the species boundary”. In: *Global Ethics & Environment*. Routledge, pp. 146–157.
- (2023). *Animal Liberation Now: The Definitive Classic Renewed*. Ed. by Yuval N. Harari. New York: Harper.
- Wittgenstein, Ludwig (1953). *Philosophical Investigations*. Ed. by G. E. M. Anscombe. New York: Wiley-Blackwell.